

Lecture

The Web Server

Outline

- Available Web Servers
- Server Features
 - Document Root
 - Authentication
 - Proxy Servers
 - Caching
 - CGI Scripting
 - Application Program Interface
- Configuring a Server
- Analyzing a Server's Performance
- Server Log Files

What Does the WWW Server Do?

- Enables browser requests
- Mainly provides
 - Support for retrieving hypertext documents
 - Manages access to the Web site
 - Provides several mechanisms for executing server-side scripts
 - Common Gateway Interface (CGI)
 - Application Program Interface (API)
 - Direct Module Interfaces (SAPI)
 - provides log files and usage statistics

What to Look for in a Web Server


- Main features
 - platform they run on;
 - complete support for HTTP 1.1 / HTTP 2
 - Multithreading, load balancing
- Security features
 - ability to provide IP address restriction
 - ability to provide domain name restriction
 - Support for secure transactions: SSL
 - Ability to act as a proxy server

How Servers Handle Multiple Requests

- For each request, a complete copy of the server is made and executed
 - Initially a parent process waits for requests; for each new request a child process is spawned
- Or, a single server program handles many requests simultaneously (multithreaded)
 - the server must keep track of all requests and switch between them as needed
 - writing multithreaded programs is easier in a language that supports multiple threads, e.g. Java

Application Web Server

- An **application server** is software that typically interfaces one or more databases to convey processed data to and from a user interface such as a web browser



- An application server acts as a set of components accessible through an API
- For web applications, these components are usually performed in the same machine where the web server is running, and their main job is to support the construction of dynamic pages.
- For example, Apache Tomcat is a popular, light-weight application container, but not a full application server as it doesn't provide the services specified in the J2EE specification. The web modules include servlets and JavaServer Pages. Business logic resides in Enterprise JavaBeans (EJB).


Some Available Web & Application Servers

Publisher	Product	Platform	URL
Web Servers			
<i>apache</i>	apache (2.4)	Win32/UNIX	http://httpd.apache.org/
<i>Microsoft</i>	IIS 7-10.0	Win32	http://www.iis.net
<i>NGINX</i>	nginx (open source)	Linux, BSD, Win32	http://nginx.org
<i>NGINX (F5)</i>	NGINX Plus	Ubuntu, AWS GCP, Azure	http://www.nginx.com
Application Servers			
<i>Oracle</i>	GlassFish 4.1	Solaris/Win32/Linux	https://javaee.github.io/glassfish/ (now Open Source Java EE Reference Implementation)
<i>IBM</i>	WebSphere	Win32/UNIX/Linux	https://www.ibm.com/cloud/websphere-application-server (free for “single developer” and “small company” use) https://www.ibm.com/us-en/marketplace/java-ee-runtime
<i>Oracle</i>	WebLogic Server	Win32/UNIX/Linux	https://www.oracle.com/java/weblogic/

For a comparison of web servers see <http://www.serverwatch.com/>, click on Server Comparison Tool


Web Server Usage

- Netcraft has identified more than 1,212,000,000 sites (4/2021)
 - Apache no longer dominates Internet servers
 - NGINX on top at 36%, Apache at #2 at 26%, Microsoft at 6%
- Statistics on intranets are difficult to determine
 - <http://news.netcraft.com/archives/category/web-server-survey/>



Web Server Survey Takeaways

- From Netcraft **April 2021** and previous surveys we can make these observations
 - **Nginx, Apache, Microsoft IIS**, have each achieved their widespread popularity in the server market due to their general availability, but many large companies in the web industry choose to roll out their own solutions.
 - **Google's custom server software ("GSE")** operates on over 37 million hostnames and 2.0 million domains. Google open sourced part of its software, the **Google Servlet Engine**, at the end of 2008, but the software has not received any official public updates in 11 years.
 - **Taobao**, China's largest online marketplace and part of Alibaba Group, developed their own **fork of nginx**. The fork, known as **Tengine**, was released back to the community as open source, leading to wider adoption. Tengine now runs on 56.8 million sites, making it the fourth most popular web server by this metric, despite it having only 1.2 million domains.



Web Server Features

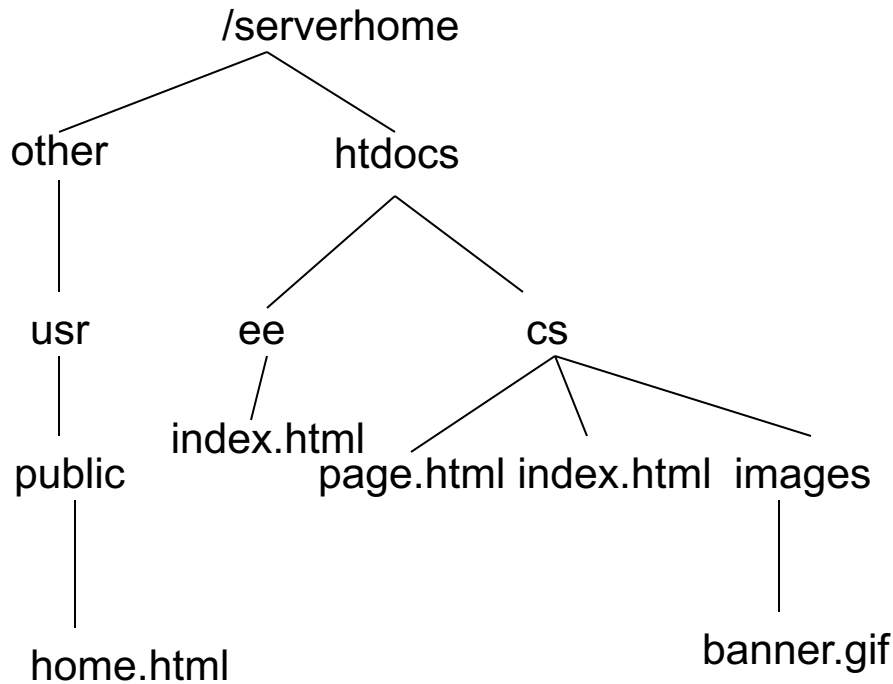
- Document Root
- Authentication
- Proxy Servers
- Caching
- CGI Scripting
- Application Program Interface

Web Server Features - Document Root

- The Web server has access to a tree of files that it can deliver
 - the files are created by content providers
 - files may contain text, images, sound, video, other programs, etc.
- the document tree is organized by the web site administrator
- The root of the document tree is given to the web server when it starts

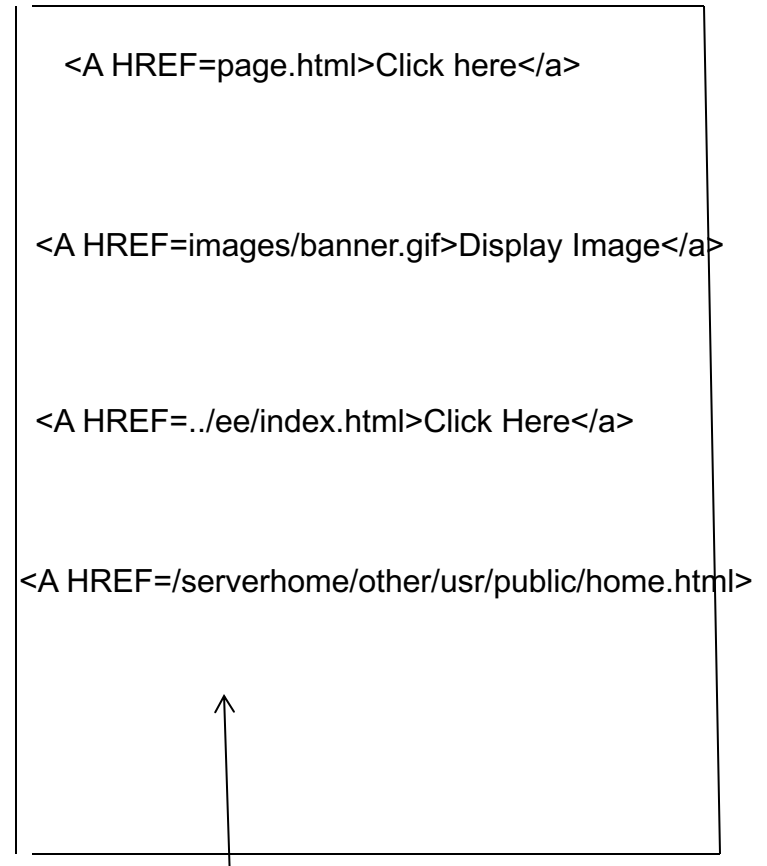
Managing the Document Tree

`http://domain/cs` is mapped to `/serverhome/htdocs/cs/index.html`



The document root is `/serverhome/htdocs`

index.html



The last link is an error

Web Server Features

Virtually Hosted Document Roots

- Hosting multiple web sites by the same web server
- It uses the host name or IP address to distinguish the document roots, e.g.

GET /index.html HTTP/1.0

Host: www.hardware.com

might return documents from /htdocs/hardware, while

GET /index.html HTTP/1.0

Host: www.antiques.com

might return documents from /htdocs/antiques

- This might be configured in, say apache by writing

<VirtualHost www.hardware.com>

ServerName www.hardware.com

DocumentRoot /htdocs/hardware

</VirtualHost>

<VirtualHost www.antiques.com>

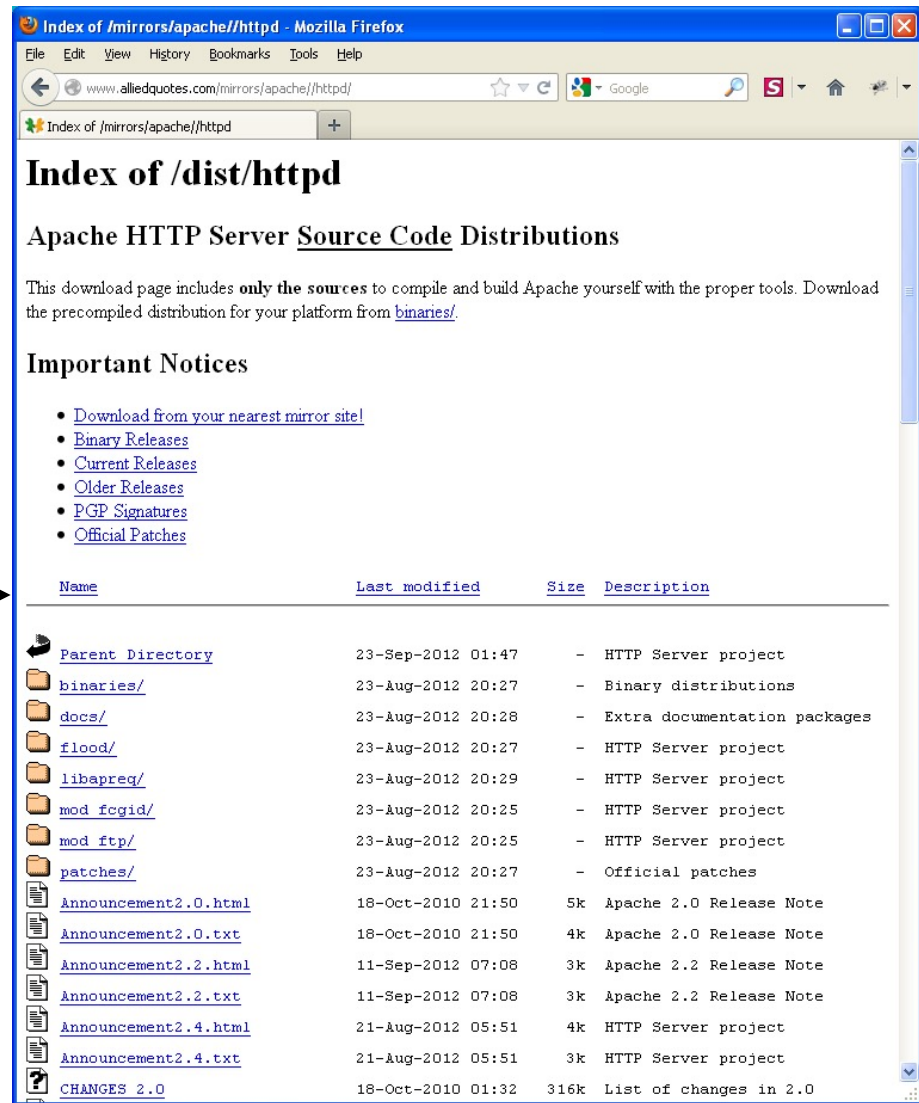
ServerName www.antiques.com

DocumentRoot /htdocs/antiques

</VirtualHost>

Web Server Features - Directory Listing

- When the URL path resolves to a directory, not to a file, a web server can be configured to return more than just a default file; to the right you see the server returning a list of files in the directory including special icons for files and folders
- Important Note:** Turn off this feature! Remove "Indexes" from **Options Indexes** in httpd.conf



Web Server Features

Basic User Authentication

- Basic authentication is supported by all HTTP servers
 - The server administrator creates secure directories accessible via password files maintained by the server
 - Client requests a resource from a secure directory; e.g., **GET /secure/test.html HTTP/1.0**
 - Server responds with authentication request to the client; e.g., **HTTP/1.0 401 Unauthorized**
 - Browser requests username and password, scrambles them, and retries the request

GET /secure/test.html HTTP/1.0

Authorization: Basic 0<V%L:EB.G-E8W)) \$J

- Server decodes name and password and checks it against its password file

Web Server Features

Creating Server-Side Applications

- Web Servers offer several mechanisms
 - Application Programming Interface (API)
 - Common Gateway Interface (CGI)
 - J2EE / .NET interfaces
 - Chrome V8 JavaScript Engine (Google)
 - Node.js
 - Direct Module Interfaces (PHP)
 - Apache offers a PHP module
- Microsoft IIS server-side supports:
 - CGI applications
 - API applications compiled as DLL
 - Active Server Pages (written in VBScript)
 - ASP.NET applications (written in C++, VB.NET, C#, or HTML/CSS/JavaScript)

Configuring a Server

- No matter which operating system or server, you will need to define the
 1. location of the server (server root)
 2. location of documents to deliver (document root)
 3. location of CGI scripts or server-side components to execute
- You may also wish to define
 - Access restrictions
 - Fancy indexing
 - Sys admin e-mail
 - Other features

Port Numbers

- A typical computer will have thousands of ports
- System ports: 0 - 1023, are reserved for certain functions, typically assigned by IANA (Internet Assigned Number Authority), <http://www.iana.org/>
- Well-known TCP / UDP ports:

http://en.wikipedia.org/wiki/List_of_TCP_and_UDP_port_numbers#Well-known_ports

- programs running on these ports must be started under root

- TCP/UDP ports used by Apple products:

<http://support.apple.com/kb/HT6175>

- Dynamic ports 49152 - 65535, ports that are freely available (check OS X possible conflicts)

Some Well-Known Port Numbers

Port	Protocol	Purpose
20	ftp	ftp data connection
21	ftp	ftp control connection
23	telnet	telnet session
25	SMTP	simple mail transfer protocol
53	DNS	domain name service
70	gopher	gopher protocol
79	finger	finger protocol
80	http	hypertext transfer protocol
88	kerberos	kerberos authentication protocol
110	pop3	post office protocol
113	ident	remote identity service
119	nntp	network news transfer protocol
143	IMAP	email access
161	snmp	simple network management proto
162	snmp	snmp traps
194	irc	internet relay chat
280	http mgmt	http management
389	ldap	lightweight directory access pr

Standard Port Numbers (cont'd)

Port	Protocol	Purpose
427	srvloc	server location protocol
443	https	secure http (SSL)
465	smtps	secure email (SSL)
535	iiop	internet inter-orb protocol
551	cybercash	secure money transaction
563	snews	secure news (SSL)
614	ssl shell	ssl shell
636	ldaps	secure ldap (SSL)
989	ftps	secure ftp data connection
990	ftps	secure ftp control connection
992	telnets	secure telnet (SSL)
993	imaps	secure imap (SSL)
994	ircs	secure irc (SSL)
995	pop3s	secure pop3 (ssl)
1080	socks	socks protocol, circuit level

The Apache Web Server

Apache Web Server

- Apache is based on code and ideas developed at NCSA in httpd version 1.3 (early 1995)
- The name “Apache” comes from the fact that it is A set of **PAtCHES** to **NCSA’s httpd**
- Apache is free and is downloadable from <http://httpd.apache.org>
- Apache included in every version of macOS:
<https://wpbeaches.com/install-apache-mysql-php-on-macos-11-big-sur-and-earlier-macos/>
- Using Apache in Windows 10
<http://dalaris.com/how-to-install-and-run-apache-web-server-on-windows-10/>
- **Apache HTTP Server version 2.4.43** or newer is required in order to operate a **TLS 1.3** web server with **OpenSSL 1.1.1**.

Apache Home Page

APACHE
HTTP SERVER PROJECT

APACHECON
September 21-23, 2021
www.apachecon.com

Essentials

- [Download!](#)
- [About](#)
- [License](#)
- [FAQ](#)
- [Security Reports](#)

Source Repositories

- [General Information](#)
- [Trunk](#)
- [2.4](#)

Documentation

- [Version 2.4](#)
- [Trunk \(dev\)](#)
- [Wiki](#)

Get Involved

- [Mailing Lists](#)
- [Bug Reports](#)
- [Developer Info](#)
- [User Support](#)

Subprojects

- [Docs](#)
- [Test](#)
- [Flood](#)
- [libapreq](#)
- [Modules](#)
- [mod_fcgid](#)
- [mod_lua](#)

Related Projects

- [Apache Traffic Server](#)
- [Apache Traffic Control](#)
- [Tomcat](#)
- [APR](#)
- [mod_perl](#)

Miscellaneous

- [Contributors](#)
- [Thanks!](#)
- [Sponsorship](#)

The Number One HTTP Server On The Internet

The Apache HTTP Server Project is an effort to develop and maintain an open-source HTTP server for modern operating systems including UNIX and Windows. The goal of this project is to provide a secure, efficient and extensible server that provides HTTP services in sync with the current HTTP standards.

The Apache HTTP Server ("httpd") was launched in 1995 and it has been the most popular web server on the Internet since April 1996. It has celebrated its 25th birthday as a project in February 2020.

The Apache HTTP Server is a project of [The Apache Software Foundation](#).

Apache httpd 2.4.46 Released 2020-08-07

The Apache Software Foundation and the Apache HTTP Server Project are pleased to [announce](#) the release of version 2.4.46 of the Apache HTTP Server ("httpd").

This latest release from the 2.4.x stable branch represents the best available version of Apache HTTP Server.

Apache HTTP Server version 2.4.43 or newer is required in order to operate a TLS 1.3 web server with OpenSSL 1.1.1.

[Download](#) | [ChangeLog for 2.4.46](#) | [Complete ChangeLog for 2.4](#) | [New Features in httpd 2.4](#)

Apache httpd 2.2 End-of-Life 2018-01-01

As previously announced, the Apache HTTP Server Project has discontinued all development and patch review of the 2.2.x series of releases.

The Apache HTTP Server Project had long committed to provide maintenance releases of the 2.2.x flavor through June of 2017. The final release 2.2.34 was published in July 2017, and no further evaluation of bug reports or security risks will be considered or published for 2.2.x releases.

Want to try out the Apache HTTP Server?

Great! We have updated our [download page](#) in an effort to better utilize our mirrors. We hope that by making it easier to use our mirrors, we will be able to provide a better download experience.

Please ensure that you [verify](#) your downloads using GPG or MD5 signatures.

Want to contribute to the Apache HTTP Server?

Awesome! Have a look at our current 'Help Wanted' listings then:

Title	Languages	Difficulty	Created
Rewrite mod_lua cache handling code	c	Intermediate	5/28/2018, 10:23:20 AM
Learn and work on httpd's internals	c	Advanced	1/2/2018, 9:19:40 AM
Websocket documentation	english,xml,lua,c	Intermediate	2/19/2016, 9:31:35 AM
Improve the Request Processing guide	xml,html,c	Advanced	2/14/2016, 3:16:11 AM
Traffic shaping with mod_lua	lua,c	Advanced	2/14/2016, 1:09:50 AM
Make a better guide on request hooks	xml,html,c	Advanced	2/8/2016, 1:46:07 AM
Write a better cache guide	xml,html	Expert	2/6/2016, 1:30:59 PM

Powered by [Help Wanted](#) - a task directory for Apache projects

Copyright © 1997-2021 The Apache Software Foundation.
Apache HTTP Server, Apache, and the Apache feather logo are trademarks of The Apache Software Foundation.

Available for
many
platforms:
Win64 (Win7-
10), Win32,
Solaris,
Linux,
macOS,
etc.

Apache httpd server

- The Apache httpd server
 - is a powerful, flexible, HTTP 1.1 & HTTP/2 compliant web server
 - implements the latest protocols, including **HTTP/2**
 - is highly configurable and extensible with third-party modules
 - can be customized by writing 'modules' using the Apache module API
 - provides full source code and comes with an unrestrictive license
 - runs on Windows, Netware, Mac OS X, Linux, and most versions of Unix, as well as several other operating systems
 - Latest release 2.4.46

Configuring Apache

- server functionality is available through modules which are either built-into or loaded into the server
- server instructions
 - apache reads its run-time configuration instructions from text files
- no GUI available (visual start/stop available on Windows)
- 182 configuration directives in the basic package
- On UNIX, Apache starts several processes to implement the server and handle requests
- On Windows there are only two processes, a parent and a child that handles all requests. Within the child all requests are handled by separate “threads”
- See <http://httpd.apache.org/docs/2.4/configuring.html>

How does Apache Work?

- Apache runs under a multitasking operating system, e.g., UNIX, macOS, Windows
- the binary is called httpd
- Key directories include:
 - conf contains configuration files
 - htdocs contains html files
 - logs contains logging data
 - cgi-bin contains executable scripts
 - icons set of *.gif files used e.g. in directory listings
 - src contains actual source code if you download the source distribution
- Apache idles, listening to an IP address and port

Configuring Apache - conf Directory

- In the conf directory there is a file provided with model configuration parameters

- httpd.conf

- This file must be edited

- Partial contents of the mime.types file

text/html	html htm
text/plain	txt
text/richtext	rtx
video/mpeg	mpeg mpg mpe
video/quicktime	qt mov
video/x-msvideo	avi
x-world/x-vrml	wrl vrml

Apache Settings - httpd.conf

- Directives are keywords followed by their value, e.g.

```
Listen 9637
```

```
ServerAdmin horowitz@usc.edu
```

```
ServerRoot /home/scf-03/csci571/WebServer/apache_2.2
```

```
ErrorLog logs/error_log
```

```
DocumentRoot "/home/scf-17/csci571c/apache2/htdocs"
```

```
ScriptAlias /cgi-bin/ "/home/scf-17/csci571c/apache2/cgi-bin/"
```

- the above are ones that *must* be set

Some Apache Settings - httpd.conf

Timeout: The number of seconds before receives
and sends time out

Timeout 300

KeepAlive: Whether or not to allow persistent
connections (more than one request per
connection). Set to "Off" to deactivate.

KeepAlive On

MaxKeepAliveRequests: The maximum number of
requests to allow during a persistent connection.
Set to 0 to allow an unlimited amount. They
recomend you leave this number high, for maximum
performance.

MaxKeepAliveRequests 100

KeepAliveTimeout: Number of seconds to wait for
the next request

KeepAliveTimeout 15

Apache Settings - httpd.conf

Server-pool size regulation. Apache dynamically adapts to the load it sees --- it tries to maintain enough server processes to handle the current load, plus a few spare servers to handle transient load spikes. It does this by periodically checking how many servers are waiting for a request. If there are fewer than MinSpareServers, it creates a new spare. If there are more than MaxSpareServers, some of the spares die off.

MinSpareServers 2

MaxSpareServers 2

StartServers 2 #number of servers to start

MaxClients 2

Apache Settings - httpd.conf

DocumentRoot: The directory out of which you will serve your documents.

DocumentRoot /home/cscixxx/WebServer/apache/htdocs

UserDir: The name of the directory which is appended onto a user's home directory if a ~user request is received.

UserDir public_html

DirectoryIndex: Name of the file or files to use as a pre-written HTML

DirectoryIndex index.html

FancyIndexing is whether you want fancy directory indexing or standard

IndexOptions FancyIndexing

Apache Settings - Scripting

```
# ScriptAlias: This controls which directories
  contain server scripts. Format: ScriptAlias
  fakename realname
ScriptAlias /cgi-bin/
  /home/cscixxx/WebServer/apache/cgi-bin/
# If you want to use server side includes, or CGI
  outside ScriptAliased directories, uncomment the
  following lines.
# AddType allows you to tweak mime.types without
  actually editing it, or to make certain files to be
  certain types. Format: AddType type/subtype ext1
# To use CGI scripts:
#AddHandler cgi-script .cgi
# To use server-parsed HTML files
#AddType text/html .shtml
#AddHandler server-parsed .shtml
```


Authentication in Apache

- There are two methods for controlling access to directories
 1. the file `access.conf` in the `conf/` directory can be used
 - placing access directives in a `<Directory>` block is the preferred method
 2. per-directory access rules can be set by a file placed in a specific directory
 - the name of the file is set by the directive `AccessFileName`
 - `.htaccess` is the default name
 - however, using `.htaccess` slows down the server;
 - NOTE: initially `.htaccess` is turned OFF
- When the server attempts to retrieve a document, it looks for an access control file in the directory or the parent directory.
- The file it looks for is set by the directive `AccessFileName`

Some Sectioning Directives

- The server configuration files use sectioning directives to define special access rights within a section or sub-directory
- Example: `<Directory>` specifies directory to which directives apply
- Options controls which server features are available
 - None
 - All
 - FollowSymLinks follow symbolic links
 - ExecCGI execution of CGI scripts allowed
 - Includes server-side includes are enabled
 - Indexes will automatically return a list of files in the directory

Limit Sectioning Directive

- `<Limit>` controls which clients can access a directory; directives within `LIMIT` include:
 - order in which `deny` and `allow` are evaluated
 - `deny from host1, host2, ...`
 - `allow from host1, host2, ...`
 - `require named-users or group-users or AuthUserFile`
 - `referer` allows access only from this directory
 - `satisfy all or any`

Using LIMIT

- Host Filtering is used to limit document trees to certain machines
- Example 1: to limit access to the cscixxx public_html documents to USC only

```
<Directory /usr/~cscixxx/public_html/>
```

```
<Limit GET>
```

```
order deny,allow
```

```
deny from all
```

```
allow from .usc.edu
```

```
</Limit>
```

```
</Directory>
```

- If someone tries to access documents in this directory from outside of usc they get a 403 Forbidden message

Using LIMIT (cont'd)

- Example 2: to limit documents so USC people CANNOT access them

```
<Directory /usr/~cscixxx/public_html/>
```

```
<Limit GET>
```

```
order allow,deny
```

```
allow from all
```

```
deny from from .usc.edu
```

```
</Limit>
```

```
</Directory>
```

Using LIMIT (cont'd)

- Example 3: a directive used to limit access to only USC and ISI domains

```
<limit GET>
```

```
order deny, allow
```

```
deny from all
```

```
allow from 128.125
```

```
allow from 128.9
```

```
</Limit>
```

Using LIMIT (cont'd)

- Example 4: Suppose you want to restrict files in a directory called secure/ to the user named student1 and password XXXYYY
- Step 1. create a file called .htaccess in directory secure/ that contains these 7 lines:

```
AuthUserFile /otherdir/.htpasswd      must be a full
    UNIX pathname
AuthGroupFile /dev/null                there is no group file
AuthName ByPassword                   any name is OK here
AuthType Basic                         Basic or Digest is OK
<Limit GET>
require user student1
</Limit>
```

Using LIMIT (cont'd)

- Step 2. Next create the password file `/otherdir/.htpasswd` using the `htpasswd` program distributed with apache.
- Type
`htpasswd -c /otherdir/.htpasswd student1`
and enter the password twice.
- You are done.
- Step 3. To add other names and passwords run `htpasswd` without the `-c` option.

Using LIMIT (cont'd)

- To generalize this example to include several students,
- Step 4. create a group file called /otherdir/.htgroup which looks like:

```
my-users: student1 student2 student3 student4
```

Alter the .htaccess file to look like

```
AuthUserFile /otherdir/.htpasswd
```

```
AuthGroupFile /otherdir/.htgroup
```

Now it points
to the group file

```
AuthName ByPassword
```

```
AuthType Basic
```

```
<Limit GET>
```

```
require group my-users
```

```
</Limit>
```

Contents of .htaccess

- The .htaccess file contains pointers to the users file, the group file, and defines the name and type of authorization

```
nunki.usc.edu(10): more .htaccess
```

```
AuthUserFile /home-scf-03/cscixxx/pwd/.htpasswd
```

```
AuthGroupFile /home-scf-03/cscixxx/pwd/.htgroup
```

```
AuthName OnlineCourse
```

```
AuthType Basic
```

Contents of .htpasswd

- This file contains a set of names and passwords

nunki.usc.edu(13): more .htpasswd

admin:kjpokufbLdrXU

chungw:2Ds5QooEqT1HM

csci351:ykGZC.bMvV8SY

csci351notes:7814Vi6oJtq9c

csci565:t.OSxBunNhigg

csci571:NUKyNCCLvgLH.

nunki.usc.edu(13): more .htgroup

admin: admin csci571 csci565

student: admin csci571 csci565

Adding a Smart Error Page

- When someone requests a page or directory that does not exist it is better to return an error page than a 404 Not Found message
- For an example try: `http://www.usc.edu/x.html`
- Look at `/conf/httpd.conf`
- find the section that describes the properties of your domain, e.g.

```
<VirtualHost www.test.com>
ServerAdmin webmaster@test.com
DocumentRoot /web/test
ServerName www.test.com
ErrorLog /log/test/error_log
transferLog /log/test/access_log
AgentLog /log/test/agent_log
RefererLog /log/test/refer_log
</VirtualHost>
```

Adding a Smart Error Page

- Insert the line
`ErrorDocument 404 /error-page.html`
- place the error page at
`DocumentRoot/ErrorDocument` or in this case
`/web/test/error-page.html`

Sample contents of `error-page.html`:

```
<HTML><HEAD><TITLE>Test Error Page</TITLE></HEAD>
<BODY>Error. You have requested a page that
    cannot be found. You may wish to return to the
    <A HREF=http://www.test.com/>home page by
    clicking here.</A>
</BODY></HTML>
```

Apache Approaches to Server-Side Processing

- CGI - the standard way to invoke an external process, e.g. perl & perl.exe
- Fast CGI, Java Servlets - Apache passes a request to another running server which executes the request and returns the result
 - see mod_fastcgi at <http://www.fastcgi.com>
- Interpreters embedded in the server itself. There are roughly two categories
 - Modules that answer or modify requests directly, e.g. mod_perl (Perl)
 - Modules aimed to process commands embedded in HTML pages before serving it to the client. e.g. mod_php5 (PHP)

Virtual Hosting

- *Virtual hosting* is a way of setting up your server so that it can appear to be multiple Web sites at once, depending upon how it's accessed by clients
- ISPs do this a lot
- The Apache Web server software allows you to create as many such virtual site identities as your system's limits will allow;
- Example of a virtual host definition.

```
<VirtualHost 10.0.0.1>
```

```
    ServerName WWW.Mydomain.Com
```

```
    ServerAdmin Webmaster@Mydomain.Com
```

```
    DocumentRoot /home/httpd/mydomain
```

```
    CustomLog logs/mydomain_log CLF
```

```
</VirtualHost>
```

Virtual Hosting (cont'd)

- there are two ways to do virtual hosting
 - *address-based* virtual hosting uses IP addresses
 - *name-based* hosting allow a single IP address to have multiple identities, e.g.

```
NameVirtualHost 10.0.0.1                this example assigns
<VirtualHost 10.0.0.1>                  two domain names to the
    ServerName WWW.Mydomain.Com         same IP address
    ServerAlias WWW.X.Com WWW.Y.Org
    ServerAdmin Webmaster@Mydomain.Com
    DocumentRoot /home/httpd/mydomain
    CustomLog logs/mydomain_log_CLF
</VirtualHost>
<VirtualHost 10.0.0.1>
    ServerName WWW.Other.Org
    ServerAlias WWW.Z.Net WWW.A.Org
    ServerAdmin Webmaster@Other.Org
    DocumentRoot /home/httpd/other
    CustomLog logs/other_log_CLF
</VirtualHost>
```


Content Negotiation

- The client provides a set of preferences concerning the preferred attributes of a document, and the server tries to come up with the best match from the versions available
 - If there's only one version of a document, then that's what the server is going to send
- The set of all versions of a document that are available to the server is called the resource's list of *variants*.
- The ways in which they differ, such as by language or encoding, are called its *dimensions*.
- The Apache server can negotiate along the
 - content-type,
 - encoding, and
 - language dimensions

Content Negotiation

- The primary use of negotiation on the Web today is to select translations of text.
 - Typically, a work is prepared in a default language such as English, and the result put in a file named `text.html` or some such. Then translated versions are prepared, and each is put into an appropriate-named file: `text.html.da` for Danish, `text.html.fr` for French, `text.html.en-gb` for English (Great Britain dialect), and so on.
- You can either prepare a list of variants yourself, or you can allow Apache's `mod_negotiation` module to calculate the list itself by scanning the directory.
- The latter method is more common, since it is easier and automatically pick up new variants as soon as they are available, but it *does* impose a slight performance penalty due to the directory scan.

Apache as a Proxy Server

- proxy servers are frequently used in fire walled environments to provide a single point of contact between the users inside and the Web outside the firewall.
 - The users ask for a page, the browser sends the request to the proxy server, the proxy server sends the request to the origin server, and the data flows back the same way.
- A second function commonly associated with proxy servers is caching.
- To use the proxy functions in your Apache Web server, the `mod_proxy` module must be activated, either by being statically linked in (less common than formerly), or turned on at run-time with directives, e.g.

AddModule mod_proxy.c

Apache as a Proxy Server, Proxy Directives

- **ProxyRequests** - This is the 'master directive,' which enables and disables proxy operation overall. It takes a single keyword argument, which must be either On or Off
- **ProxyRemote** - tells the Apache server that certain requests are to be forwarded to yet another system rather than being handled locally, allowing for cascading of proxies
- **ProxyBlock** - allows you to selectively block the fetching of certain resources through the proxy. It takes a list of space-separated case-insensitive phrases or URI substrings which should be blocked.
 - "ProxyBlock foo" would block all requests in which "foo" appeared, in any combination of upper- and lower-case, including "http://foo.com/bar/" and "http://bar.org/zardofoo.html". You can block all proxy requests with "ProxyBlock *"

Controlling the Apache Cache

- **CacheRoot** - Defines the directory that's the top of the cache hierarchy
 - Apache uses its own naming conventions for the stuff it stores.
- **CacheSize** - Specifies the largest number of kilobytes of space you want to allocate to the proxy cache.
- **CacheDirLevels** - Controls how deep the cache directory tree is allowed to grow
- **CacheDirLength** - Used to define the number of characters used in naming each level of cache subdirectory
- **NoCache** - Identifies proxy-accessed documents that shouldn't be cached. The format is the same as for the ProxyBlock directive.

Proxy Host Example

- Usually a proxy server is set up as a completely separate virtual host from the rest of the server's operation (which is likely to be handling normal requests).
- You can do this with a `<VirtualHost>` stanza. The `<IfModule>` container will only be evaluated if the proxy module is part of the running server

```
<IfModule mod_proxy.c>
```

```
    Listen 10.0.0.1:8080
```

```
    <VirtualHost 10.0.0.1:8080>
```

```
        ProxyRequests On
```

```
        DocumentRoot /home/httpd/htdocs/proxy
```

```
    </VirtualHost>
```

```
</IfModule>
```

Analyzing Web Server Log Files

- A sample entry from a log file

Host	date/time	method	file	HTTP ver	status	bytes
marmot.usc.edu	- - [17/Feb/2009:23:21:09 -0800]	"GET	/	HTTP/1.0"	200	1316

- There are different log file formats
 - W3C Extended Log File Format
 - Date/time, server port, URL and query, protocol version, user agent and cookie, referrer, etc.
 - Microsoft IIS log file format
 - IP, user, date/time, server name, server IP, bytes received, bytes sent, protocol, method, etc.
 - Other file formats
 - Cerfnet, apache extended, CERN common, Lotus Domino, Netscape Common, etc.

Definition of Terms

- *A hit*: a request for a page, image, cgi script, etc.
- *A page view*: a request for an html page
- Unique host vs. unique visitor: aol.com is a unique host, so this statistic may not be useful; cookies are best way to track unique visitors
- *An impression*: is a view of an advertisement
- *A click-through* is when an advertisement is actually clicked on

Common Log Format

- Each visit to your site produces a record of the request in log file
- log files in Apache are typically stored in the logs directory
- Some fields include:
 - Host client hostname or IP address
 - ident remote identity, often a dash
 - usr authenticated user name
 - time date/time request is received
 - req HTTP request
 - s1 server's HTTP response status code
 - c1 Content-length
- Example entries from access log

```
marmot.usc.edu - - [17/Feb/2008:23:21:09 -0800]  
"GET / HTTP/1.0" 200 1316
```

```
128.125.229.195 - - [23/Feb/2008:10:32:09 -0800]  
"GET / HTTP/1.1" 200 1316
```

Other Entries in access Log

- ciscoserv-229-180.usc.edu - - [17/Feb/2008:23:16:51 -0800] "GET /cgi-bin/secure/admin/student.pl HTTP/1.1" 401 362
- 128.125.229.195 - admin [23/Feb/2008:10:56:21 -0800] "POST /cgi-bin/secure/admin/student.pl HTTP/1.1" 200 807
- 209.85.18.67 - - [27/Feb/2008:01:12:38 -0800] "GET /ftp/templates.zip HTTP/1.1" 200 30780
- nunki.usc.edu - - [28/Feb/2008:00:48:22 -0800] "GET / HTTP/1.0" 200 955
- nunki.usc.edu - - [28/Feb/2008:00:48:26 -0800] "GET /cgi-bin/Count.cgi?df=count.dat HTTP/1.0" 200 207

access_log Contents

- Here are some lines from a single user session

```
lv1-sun702.usc.edu - - [09/Feb/2008:19:15:36
-0800] "GET /newsdigest/index.html HTTP/1.0"
200 3630

lv1-sun702.usc.edu - - [09/Feb/2008:19:15:37
-0800] "GET / HTTP/1.0" 200 955

lv1-sun702.usc.edu - - [09/Feb/2008:19:15:39
-0800] "GET /newsdigest/index.html HTTP/1.0"
200 3630
```
- A *user session* includes all hits and requests made by a user terminated by “n” minutes of inactivity

Sample error_log Contents

- Here are some lines from the error_log file

- **The file was not found**

- ```
[Mon Jan 29 11:47:36 2008] httpd: access to /home/scf-13/csci665/www/pleader/horohome.html failed for internet-gw.watson.ibm.com, reason: No file matching URL: /pleader/horohome.html from -
```

- **authorization entered is incorrect**

- ```
[Thu Sep 24 10:17:41 2008] [error] [client 128.125.50.47] user csci571 not found: /cgi-bin/secure/test-cgi
```

- **Perl script fails**

- ```
[Sun Sep 30 01:15:01 2008] [error] [client 71.140.65.229] Premature end of script headers: /auto/home-scf-22/csci571/WebServer/Apache_1.3.36/cgi-bin/cgipm/test6.pl
```

## More Examples of Error Log File Entries

- Undefined subroutine &main::Export\_to\_Excel called at /home/scf-03/csci571/WebServer/apache\_1.2.5/cgi-bin/admin/student\_db.pl line 75.
- [Tue Dec 29 02:07:47 2008] access to /auto/home-scf-03/csci571/WebServer/apache\_1.2.5/htdocs/robots.txt failed for 193.189.227.21, reason: File does not exist
- [Tue Dec 29 20:11:22 2008] send directory lost connection to client 128.125.223.237

# Performance and Security Considerations

- Don't log more than needed
- Look for 401 (unauthorized) and 403 forbidden status codes
- Look at the response times
  - Which parts of the web site have the longest response times?
  - Where are the bottlenecks
- Look for failed page hits (404, 5xx)
- Look for cached hits
- Look for browsers and platforms of clients

# Tools to Analyze Log Files

- Some commercially available tools for IIS
  - **WebTrends**  
<http://www.webtrends.com>  
See also  
[http://cs-server.usc.edu:45678/resources/web\\_servers/webtrends/report2.html](http://cs-server.usc.edu:45678/resources/web_servers/webtrends/report2.html)
  - Marketware, webmanager, MS site server, LiveStats
- For any web server
  - **Google Analytics**. See:  
<http://www.google.com/analytics/>

# Google Analytics

